

Decision Trees & Random-Forests using Lending Club Data ¶

Brief

For this project we will be exploring publicly available data from www.LendingClub.com (<http://www.LendingClub.com>). Lending Club connects people who need money (borrowers) with people who have money (investors). Hopefully, as an investor you would want to invest in people who showed a profile of having a high probability of paying you back. We will try to create a model that will help predict this.

Lending club had a very interesting year in 2016, this data is from them before they even went public. The lending data is from 2007-2010 and trying to classify and predict whether or not the borrower paid back their loan in full.

Here are what the columns represent:

- credit.policy: 1 if the customer meets the credit underwriting criteria of LendingClub.com, and 0 otherwise.
- purpose: The purpose of the loan (takes values "credit_card", "debt_consolidation", "educational", "major_purchase", "small_business", and "all_other").
- int.rate: The interest rate of the loan, as a proportion (a rate of 11% would be stored as 0.11). Borrowers judged by LendingClub.com to be more risky are assigned higher interest rates.
- installment: The monthly installments owed by the borrower if the loan is funded.
- log.annual.inc: The natural log of the self-reported annual income of the borrower.
- dti: The debt-to-income ratio of the borrower (amount of debt divided by annual income).
- fico: The FICO credit score of the borrower.
- days.with.cr.line: The number of days the borrower has had a credit line.
- revol.bal: The borrower's revolving balance (amount unpaid at the end of the credit card billing cycle).
- revol.util: The borrower's revolving line utilization rate (the amount of the credit line used relative to total credit available).
- inq.last.6mths: The borrower's number of inquiries by creditors in the last 6 months.
- delinq.2yrs: The number of times the borrower had been 30+ days past due on a payment in the past 2 years.
- pub.rec: The borrower's number of derogatory public records (bankruptcy filings, tax liens, or judgments).



```
In [4]: import numpy as np
import pandas as pd
```

```
In [5]: import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [6]: loans= pd.read_csv('loan_data.csv')
```

```
In [17]: loans.info ()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
credit.policy      9578 non-null int64
purpose           9578 non-null object
int.rate          9578 non-null float64
installment       9578 non-null float64
log.annual.inc    9578 non-null float64
dti               9578 non-null float64
fico              9578 non-null int64
days.with.cr.line 9578 non-null float64
revol.bal         9578 non-null int64
revol.util        9578 non-null float64
inq.last.6mths    9578 non-null int64
delinq.2yrs       9578 non-null int64
pub.rec           9578 non-null int64
not.fully.paid    9578 non-null int64
dtypes: float64(6), int64(7), object(1)
memory usage: 1.0+ MB
```

In [7]: `loans.head(10)`

Out[7]:

	credit.policy	purpose	int.rate	installment	log.annual.inc	dti	fico	days.with.cr.line	revol.bal	revol.util	inq.last.6mths	delinq.2yrs
0	1	debt_consolidation	0.1189	829.10	11.350407	19.48	737	5639.958333	28854	52.1	0	0
1	1	credit_card	0.1071	228.22	11.082143	14.29	707	2760.000000	33623	76.7	0	0
2	1	debt_consolidation	0.1357	366.86	10.373491	11.63	682	4710.000000	3511	25.6	1	0
3	1	debt_consolidation	0.1008	162.34	11.350407	8.10	712	2699.958333	33667	73.2	1	0
4	1	credit_card	0.1426	102.92	11.299732	14.97	667	4066.000000	4740	39.5	0	1
5	1	credit_card	0.0788	125.13	11.904968	16.98	727	6120.041667	50807	51.0	0	0
6	1	debt_consolidation	0.1496	194.02	10.714418	4.00	667	3180.041667	3839	76.8	0	0
7	1	all_other	0.1114	131.22	11.002100	11.08	722	5116.000000	24220	68.6	0	0
8	1	home_improvement	0.1134	87.19	11.407565	17.25	682	3989.000000	69909	51.1	1	0
9	1	debt_consolidation	0.1221	84.12	10.203592	10.00	707	2730.041667	5630	23.0	1	0

In [10]: `loans.describe()`

Out[10]:

	nt.rate	installment	log.annual.inc	dti	fico	days.with.cr.line	revol.bal	revol.util	inq.last.6mths	delinq.2yrs	pub.rec
count	100000	9578.000000	9578.000000	9578.000000	9578.000000	9578.000000	9.578000e+03	9578.000000	9578.000000	9578.000000	9578.000000
mean	22640	319.089413	10.932117	12.606679	710.846314	4560.767197	1.691396e+04	46.799236	1.577469	0.163708	0.062122
std	126847	207.071301	0.614813	6.883970	37.970537	2496.930377	3.375619e+04	29.014417	2.200245	0.546215	0.262126
min	160000	15.670000	7.547502	0.000000	612.000000	178.958333	0.000000e+00	0.000000	0.000000	0.000000	0.000000
25%	103900	163.770000	10.558414	7.212500	682.000000	2820.000000	3.187000e+03	22.600000	0.000000	0.000000	0.000000
50%	22100	268.950000	10.928884	12.665000	707.000000	4139.958333	8.596000e+03	46.300000	1.000000	0.000000	0.000000
75%	40700	432.762500	11.291293	17.950000	737.000000	5730.000000	1.824950e+04	70.900000	2.000000	0.000000	0.000000
max	116400	940.140000	14.528354	29.960000	827.000000	17639.958330	1.207359e+06	119.000000	33.000000	13.000000	5.000000


Exploratory Data Analysis¶ (EDA)

- CreateING a histogram of two FICO distributions on top of each other, one for each credit.policy outcome.

In [11]:

```
plt.figure(figsize=(10,6))
loans[loans['credit.policy']==1]['fico'].hist(alpha=0.5,color='blue',
                                              bins=30,label='Credit.Policy=1')
loans[loans['credit.policy']==0]['fico'].hist(alpha=0.5,color='red',
                                              bins=30,label='Credit.Policy=0')
plt.legend()
plt.xlabel('FICO')
```


Out[11]: Text(0.5, 0, 'FICO')



```
In [8]: plt.figure(figsize=(10,6))
        loans[loans['not.fully.paid']==1]['fico'].hist(alpha=0.5,color='green',
                                                    bins=25,label='not.fully paid =1')
        loans[loans['not.fully.paid']==0]['fico'].hist(alpha=0.5,color='brown',
                                                    bins=25,label='not.fully paid=0')


        plt.legend()
        plt.xlabel('FICO')
```

Out[8]: Text(0.5, 0, 'FICO')



```
In [15]: plt.figure(figsize=(10,7))  
sns.countplot(x='purpose', hue='not.fully.paid', data=loans,)
```


```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x1e5463b2588>
```




```
In [11]: plt.figure(figsize=(10,6))  
sns.jointplot(x='fico',y='int.rate',data=loans, color='purple')
```

```
Out[11]: <seaborn.axisgrid.JointGrid at 0x25017445070>
```


```
<Figure size 720x432 with 0 Axes>
```



```
In [19]: plt.figure(figsize=(10,7))
sns.lmplot(y='int.rate',x='fico', data=loans,hue='credit.policy', col='not.fully.paid')
```

Out[19]: <seaborn.axisgrid.FacetGrid at 0x256aad58a88>

<Figure size 720x504 with 0 Axes>



Setting up the Data

Random Forest Classification Model!

The loans data has an object columns will use `pd.get_dummies` to make it an int

```
In [113]: cat_feats=['purpose']
```

```
In [114]: final_data = pd.get_dummies(loans,columns=cat_feats,drop_first=True )
```

```
In [29]: final_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 19 columns):
credit.policy      9578 non-null int64
int.rate           9578 non-null float64
installment        9578 non-null float64
log.annual.inc     9578 non-null float64
dti                9578 non-null float64
fico               9578 non-null int64
days.with.cr.line 9578 non-null float64
revol.bal          9578 non-null int64
revol.util          9578 non-null float64
inq.last.6mths     9578 non-null int64
delinq.2yrs        9578 non-null int64
pub.rec            9578 non-null int64
not.fully.paid     9578 non-null int64
purpose_credit_card 9578 non-null uint8
purpose_debt_consolidation 9578 non-null uint8
purpose_educational 9578 non-null uint8
purpose_home_improvement 9578 non-null uint8
purpose_major_purchase 9578 non-null uint8
purpose_small_business 9578 non-null uint8
dtypes: float64(6), int64(7), uint8(6)
memory usage: 1.0 MB
```

```
In [28]: final_data.head()
```

```
Out[28]:
```

	credit.policy	int.rate	installment	log.annual.inc	dti	fico	days.with.cr.line	revol.bal	revol.util	inq.last.6mths	delinq.2yrs	pub.rec	not.fully.paid
0	1	0.1189	829.10	11.350407	19.48	737	5639.958333	28854	52.1	0	0	0	0
1	1	0.1071	228.22	11.082143	14.29	707	2760.000000	33623	76.7	0	0	0	0
2	1	0.1357	366.86	10.373491	11.63	682	4710.000000	3511	25.6	1	0	0	0
3	1	0.1008	162.34	11.350407	8.10	712	2699.958333	33667	73.2	1	0	0	0
4	1	0.1426	102.92	11.299732	14.97	667	4066.000000	4740	39.5	0	1	0	0

```
In [115]: from sklearn.model_selection import train_test_split
```

```
In [116]: x= final_data.drop('not.fully.paid',axis = 1)
y= final_data['not.fully.paid']
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=101)
```

Training a Decision Tree Model¶

Let's start by Training a single decision tree first!

Import DecisionTreeClassifier

```
In [117]: from sklearn.tree import DecisionTreeClassifier
```

```
In [118]: dtree=DecisionTreeClassifier ()
```

```
In [119]: dtree.fit(x_train,y_train)
```

```
Out[119]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False,
                                random_state=None, splitter='best')
```

Predictions and Evaluation of Decision Tree

Creating a classification report and a confusion matrix

```
In [140]: predictions = dtree.predict(x_test)
```

```
In [141]: from sklearn.metrics import classification_report, confusion_matrix
```

```
In [142]: print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.85	0.82	0.84	2431
1	0.19	0.23	0.21	443
accuracy			0.73	2874
macro avg	0.52	0.53	0.52	2874
weighted avg	0.75	0.73	0.74	2874

```
In [146]: print (confusion_matrix(y_test, predictions))
```

```
[[1996  435]
 [ 340  103]]
```

Training the Random Forest model

train our model!

RandomForestClassifier class and fit it to our training data

```
In [147]: from sklearn.ensemble import RandomForestClassifier
```

```
In [149]: rfc = RandomForestClassifier(n_estimators=600)
```

```
In [151]: rfc.fit(x_train,y_train)
```

```
Out[151]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',  
                                max_depth=None, max_features='auto', max_leaf_nodes=None,  
                                min_impurity_decrease=0.0, min_impurity_split=None,  
                                min_samples_leaf=1, min_samples_split=2,  
                                min_weight_fraction_leaf=0.0, n_estimators=600,  
                                n_jobs=None, oob_score=False, random_state=None,  
                                verbose=0, warm_start=False)
```

Predictions and Evaluation

predicting off the y_test values and evaluate the model

Predict the class of not.fully.paid for the X_test data

```
In [152]: prediction = rfc.predict(x_test)
```

```
In [154]: from sklearn.metrics import classification_report, confusion_matrix
```

```
In [155]: print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.85	0.82	0.84	2431
1	0.19	0.23	0.21	443
accuracy			0.73	2874
macro avg	0.52	0.53	0.52	2874
weighted avg	0.75	0.73	0.74	2874

```
In [156]: print(confusion_matrix(y_test,predictions))
```

```
[[1996  435]
 [ 340  103]]
```

Conclusion: from the model above, it shows people are paying back their loans after all Criteria of lending Club are met