


**The project is from university Buckingham from their Business Analyst department, it was undertaking for practices purpose and development of my skill set**

The task is focus on Python and R scripting for data analysis. I opted to used Python to enhance my skills and delve deeper into data analysis to the best of my ability.

The project comprised two main sections. The first section, labeled as Section A, centered on analyzing on suicide rates. Specifically, the objective was to identify vulnerable age groups. To achieve this,I examined the suicide rates across various countrie, with the year ranging from 1985 to 2016, this dataset was gather from WHO and named as WHO\_Suicide\_Data.csv.

## **Section B**

This is to demonstrate data acquiotion (Scraping, cleaning and exploration).



In [10]: *# import the libraries*

```
import numpy as np
import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
```

```
%matplotlib inline
import missingno as msno
```

## Section A

### Importaing the data set

```
In [13]: who=pd.read_csv('WHO_Suicide_Data_3.csv')
```

```
In [14]: #checking the head of the data
who.head()
```

```
Out[14]:
```

	country	year	sex	age	suicides_no	population	gdp_for_year (\$)	Unnamed: 7	Unnamed: 8
0	Albania	1987	male	15-24 years	21	312900	2,156,624,900	NaN	NaN
1	Albania	1987	male	35-54 years	16	308000	2,156,624,900	NaN	NaN
2	Albania	1987	female	15-24 years	14	289700	2,156,624,900	NaN	NaN
3	Albania	1987	male	75+ years	1	21800	2,156,624,900	NaN	NaN
4	Albania	1987	male	25-34 years	9	274300	2,156,624,900	NaN	NaN

```
In [15]: who.isnull().any()
```

```
Out[15]:
```

country	False
year	False
sex	False
age	False
suicides_no	True
population	False
gdp_for_year (\$)	False
Unnamed: 7	True
Unnamed: 8	True
dtype: bool	

```
In [16]: # Drop the 'Unnamed: 8' column

who.drop(columns=['Unnamed: 7','Unnamed: 8'], inplace=True)
```

```
In [17]: who.head()
```

```
Out[17]:
```

	country	year	sex	age	suicides_no	population	gdp_for_year (\$)
0	Albania	1987	male	15-24 years	21	312900	2,156,624,900
1	Albania	1987	male	35-54 years	16	308000	2,156,624,900
2	Albania	1987	female	15-24 years	14	289700	2,156,624,900
3	Albania	1987	male	75+ years	1	21800	2,156,624,900
4	Albania	1987	male	25-34 years	9	274300	2,156,624,900

## Checking for types

```
In [18]: who.dtypes
```

```
Out[18]: country          object
year              int64
sex              object
age              object
suicides_no      object
population        int64
gdp_for_year ($)  object
dtype: object
```

## Checking the column of the data frame

```
In [22]: print(who.columns)
```

```
Index(['country', 'year', 'sex', 'age', 'suicides_no', 'population',
      'gdp_for_year ($)'],
      dtype='object')
```

## Converting the the data types

```
In [58]: # Clean and convert 'gdp_for_year ($)' column to numeric
def clean_and_convert(value):
    cleaned_value = ''.join(c for c in str(value) if c.isnumeric() or c == '.')
    return float(cleaned_value)
```

```

who[' gdp_for_year '] = who[' gdp_for_year '].apply(clean_and_convert)

# Set display option to avoid scientific notation
pd.options.display.float_format = '{:.2f}'.format

# Print the first few rows to verify the changes
print(who.head())

```

	country	year	sex	age	suicides_no	population \
0	Albania	1987	male	15-24 years	21.00	312900
1	Albania	1987	male	35-54 years	16.00	308000
2	Albania	1987	female	15-24 years	14.00	289700
3	Albania	1987	male	75+ years	1.00	21800
4	Albania	1987	male	25-34 years	9.00	274300

	gdp_for_year	suicides/100k	generation
0	2156624900.00	6.71	Millennials
1	2156624900.00	5.19	Millennials
2	2156624900.00	4.83	Millennials
3	2156624900.00	4.59	Millennials
4	2156624900.00	3.28	Millennials

## Checking the suicide\_no

```

In [24]: non_numeric_mask = pd.to_numeric(who['suicides_no'], errors='coerce').isna()
problematic_values = who.loc[non_numeric_mask, 'suicides_no'].unique()

print(problematic_values)

[nan 'Null' 'Unknown']

```

## Converting the suicides\_no column and gdp\_for\_year from Object

```

In [59]: # Convert non-numeric values to NaN
who['suicides_no'] = pd.to_numeric(who['suicides_no'], errors='coerce')

# Convert 'gdp_for_year' column to float
who[' gdp_for_year '] = who[' gdp_for_year '].astype(float)

# Print data types to confirm the changes
print(who.dtypes)

```

```
country      object
year         int64
sex          object
age          object
suicides_no  float64
population   int64
gdp_for_year float64
suicides/100k float64
generation   object
dtype: object
```

## Confirming the dtypes

```
In [26]: who.dtypes
```

```
Out[26]: country      object
year         int64
sex          object
age          object
suicides_no  float64
population   int64
gdp_for_year ($)  object
dtype: object
```

```
In [27]: non_numeric_mask = pd.to_numeric(who['suicides_no'], errors='coerce').isna()
remaining_non_numeric_values = who.loc[non_numeric_mask, 'suicides_no'].unique()
print(remaining_non_numeric_values)
```

```
[nan]
```

```
In [28]: unique_suicides_values = who['suicides_no'].unique()
print(unique_suicides_values)
```

```
[ 21.  16.  14. ... 11634. 4359. 2872.]
```

```
In [ ]:
```

```
In [29]: who['suicides_no'] = who['suicides_no'].replace(['Null', 'Unknown'], np.nan)
```

```
In [30]: who.head()
```

```
Out[30]:
```

	country	year	sex	age	suicides_no	population	gdp_for_year (\$)
0	Albania	1987	male	15-24 years	21.0	312900	2,156,624,900
1	Albania	1987	male	35-54 years	16.0	308000	2,156,624,900
2	Albania	1987	female	15-24 years	14.0	289700	2,156,624,900
3	Albania	1987	male	75+ years	1.0	21800	2,156,624,900
4	Albania	1987	male	25-34 years	9.0	274300	2,156,624,900

## Replacing the dollar sign from the head

```
In [31]: # Replace the dollar sign in the column name
who.columns = who.columns.str.replace('gdp_for_year \(\$\)', 'gdp_for_year')
```

C:\Users\Autoke Pro\AppData\Local\Temp\ipykernel\_2532\1785478471.py:2: FutureWarning: The default value of regex will change from True to False in a future version.

```
who.columns = who.columns.str.replace('gdp_for_year \(\$\)', 'gdp_for_year')
```

```
In [32]: who.head()
```

```
Out[32]:
```

	country	year	sex	age	suicides_no	population	gdp_for_year
0	Albania	1987	male	15-24 years	21.0	312900	2,156,624,900
1	Albania	1987	male	35-54 years	16.0	308000	2,156,624,900
2	Albania	1987	female	15-24 years	14.0	289700	2,156,624,900
3	Albania	1987	male	75+ years	1.0	21800	2,156,624,900
4	Albania	1987	male	25-34 years	9.0	274300	2,156,624,900

## Dropping any duplicate in the data set

```
In [33]: # Drop duplicate rows and update the DataFrame
who = who.drop_duplicates()
```

## Checking for Duplicate

```
In [34]: # Check for duplicate rows
duplicate_rows = who[who.duplicated()]

num_duplicate_rows = who.duplicated().sum()
print(f"Number of duplicate rows: {num_duplicate_rows}")
```

Number of duplicate rows: 0

## Checking for missing value

```
In [35]: who.isnull().sum()
```

```
Out[35]: country          0
year          0
sex           0
age           0
suicides_no    4281
population     0
  gdp_for_year    0
dtype: int64
```


## Used another library to visualise the missing data section

The grey and white area is the section where value are missing in the data set

```
In [36]: msno.matrix (who)
```

```
Out[36]: <AxesSubplot:>
```





Finding the mean of missing value and going to use it to fill the NaN value. Note the missing value are in series, so fill forward or backward would not apply.

```
In [37]: mean_missing_value_percent = who.isnull().mean()*100
print (mean_missing_value_percent)
```

```
country      0.00000
year         0.00000
sex          0.00000
age          0.00000
suicides_no  15.38821
population   0.00000
gdp_for_year 0.00000
dtype: float64
```

Filling NaN with the Calculated mean of the suicides\_no column

```
In [38]: # Replace NaN values in 'suicides_no' column with mean value
mean_missing_value_percent = 15.32
```

```
who['suicides_no'] = who['suicides_no'].fillna(mean_missing_value_percent)
```

```
#who.interpolate()  
new_who=who
```

C:\Users\Autoke Pro\AppData\Local\Temp\ipykernel\_2532\732802329.py:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
who['suicides_no'] = who['suicides_no'].fillna(mean_missing_value_percent)
```

```
In [39]: new_who.head(12)
```

```
Out[39]:
```

	country	year	sex	age	suicides_no	population	gdp_for_year
0	Albania	1987	male	15-24 years	21.00	312900	2,156,624,900
1	Albania	1987	male	35-54 years	16.00	308000	2,156,624,900
2	Albania	1987	female	15-24 years	14.00	289700	2,156,624,900
3	Albania	1987	male	75+ years	1.00	21800	2,156,624,900
4	Albania	1987	male	25-34 years	9.00	274300	2,156,624,900
5	Albania	1987	female	75+ years	1.00	35600	2,156,624,900
6	Albania	1987	female	35-54 years	6.00	278800	2,156,624,900
7	Albania	1987	female	25-34 years	4.00	257200	2,156,624,900
8	Albania	1987	male	55-74 years	1.00	137500	2,156,624,900
9	Albania	1987	female	5-14 years	15.32	311000	2,156,624,900
10	Albania	1987	female	55-74 years	15.32	144600	2,156,624,900
11	Albania	1987	male	5-14 years	15.32	338200	2,156,624,900

## Confirming the suicides\_no column

```
In [40]: nan_mask = who['suicides_no'].isna()  
nan_values = who.loc[nan_mask, 'suicides_no']
```

```
print(nan_values)
```

```
Series([], Name: suicides_no, dtype: float64)
```

## Confirming for missing value

```
In [41]: new_who.isnull().sum()
```

```
Out[41]: country          0
year          0
sex           0
age           0
suicides_no    0
population     0
gdp_for_year   0
dtype: int64
```

## Adding a new column and it will be called 'suicides/100k'

```
In [42]: new_who['suicides/100k']=new_who['suicides_no']/ (new_who['population'] /100000)
```

C:\Users\Autoke Pro\AppData\Local\Temp\ipykernel\_2532\3575281647.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
`new_who['suicides/100k']=new_who['suicides_no']/ (new_who['population'] /100000)`

```
In [43]: new_who.head()
```

```
Out[43]:
```

	country	year	sex	age	suicides_no	population	gdp_for_year	suicides/100k
0	Albania	1987	male	15-24 years	21.0	312900	2,156,624,900	6.711409
1	Albania	1987	male	35-54 years	16.0	308000	2,156,624,900	5.194805
2	Albania	1987	female	15-24 years	14.0	289700	2,156,624,900	4.832585
3	Albania	1987	male	75+ years	1.0	21800	2,156,624,900	4.587156
4	Albania	1987	male	25-34 years	9.0	274300	2,156,624,900	3.281079

## Creating another column called 'generation'

```
In [60]: def assign_generation(year):  
    if 1883 <= year <= 1900:  
        return 'Lost Generation'  
    elif 1901 <= year <= 1927:  
        return 'G.I. Generation'  
    elif 1928 <= year <= 1945:  
        return 'Silent'  
    elif 1946 <= year <= 1964:  
        return 'Boomers'  
    elif 1965 <= year <= 1980:  
        return 'Generation X'  
    elif 1981 <= year <= 1996:  
        return 'Millennials'  
    elif 1997 <= year <= 2012:  
        return 'Generation Z'  
    elif 2013 <= year <= 2025:  
        return 'Generation A'  
    else:  
        return 'Unknown'  
  
new_who['generation'] = new_who['year'].apply(assign_generation)
```

```
In [45]: new_who.head()
```

```
Out[45]:
```

	country	year	sex	age	suicides_no	population	gdp_for_year	suicides/100k	generation
0	Albania	1987	male	15-24 years	21.0	312900	2,156,624,900	6.711409	Millennials
1	Albania	1987	male	35-54 years	16.0	308000	2,156,624,900	5.194805	Millennials
2	Albania	1987	female	15-24 years	14.0	289700	2,156,624,900	4.832585	Millennials
3	Albania	1987	male	75+ years	1.0	21800	2,156,624,900	4.587156	Millennials
4	Albania	1987	male	25-34 years	9.0	274300	2,156,624,900	3.281079	Millennials

## Crteating another column called 'gdp\_per\_capita'

```
In [61]: # Add the 'gdp_per_capita' column  
new_who['gdp_per_capita'] = new_who[' gdp_for_year '] / (new_who['population'] / 100000)
```

```
# Perform calculations on the 'gdp_per_capita' column
new_who['gdp_per_capita'] = new_who['gdp_per_capita'] / new_who['population'] * 100000
```

```
In [47]: new_who.head()
```

```
Out[47]:
```

	country	year	sex	age	suicides_no	population	gdp_for_year	suicides/100k	generation
0	Albania	1987	male	15-24 years	21.0	312900	2,156,624,900	6.711409	Millennials
1	Albania	1987	male	35-54 years	16.0	308000	2,156,624,900	5.194805	Millennials
2	Albania	1987	female	15-24 years	14.0	289700	2,156,624,900	4.832585	Millennials
3	Albania	1987	male	75+ years	1.0	21800	2,156,624,900	4.587156	Millennials
4	Albania	1987	male	25-34 years	9.0	274300	2,156,624,900	3.281079	Millennials

## Ranking the countries by total suicides rate

```
In [49]: # Assuming 'who' is the name of your DataFrame
# Group data by country and calculate total suicides
country_suicides = who.groupby('country')['suicides_no'].sum().reset_index()

# Sort the data by total suicides in descending order
country_suicides = country_suicides.sort_values(by='suicides_no', ascending=False)

# Create a Larger bar plot using Seaborn with adjusted font size
plt.figure(figsize=(12, 15)) # Increase both width and height
sns.barplot(x='suicides_no', y='country', data=country_suicides, palette='viridis')
plt.xlabel('Total Suicides', fontsize=14) # Increase font size
plt.ylabel('Country', fontsize=14) # Increase font size
plt.title('Ranks of Countries by Total Suicides rate', fontsize=16) # Increase font size


# Rotate country labels for better readability
plt.xticks(rotation=45, ha='right', fontsize=12) # Rotate labels and set alignment

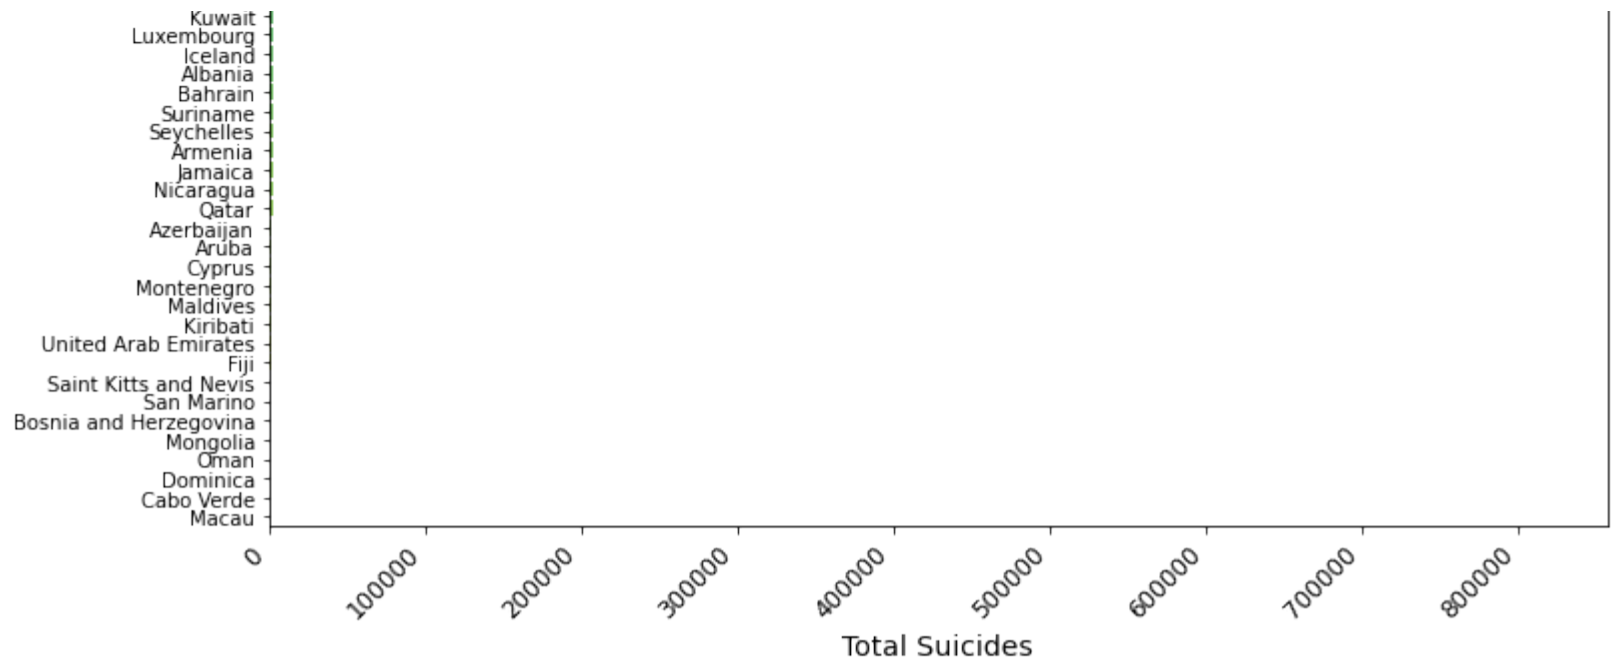
# Adjust spacing to avoid overlapping labels
plt.tight_layout()

plt.show()
```

# Ranks of Countries by Total Suicides rate

Country






## Finding the correlations between suicides, GDP per capita and population

```
In [63]: correlations = new_who [['suicides_no', 'gdp_per_capita', 'population']].corr()
print(correlations)
```

	suicides_no	gdp_per_capita	population
suicides_no	1.00	-0.01	0.55
gdp_per_capita	-0.01	1.00	-0.02
population	0.55	-0.02	1.00

```
In [64]: hep = new_who[['suicides_no', 'gdp_per_capita', 'population']]

plt.figure(figsize=(10, 8))
sns.heatmap(hep.corr().round(2), annot=True, vmin=-1, vmax=1, cmap='vlag')
plt.show()
```



My conclusion from the correlation graph, From my point of view, there is a very high rate in suicides\_no


Use appropriate visual notation to visualise total suicides over years. Describe your findings

```
In [53]: # Group the data by year and calculate the total suicides for each year
total_suicides_by_year = new_who.groupby('year')['suicides_no'].sum()

# Create a line plot to visualize total suicides over the years
plt.figure(figsize=(12, 6))
plt.plot(total_suicides_by_year.index, total_suicides_by_year.values, marker='o', linestyle='--')
plt.title('Total Suicides Over the Years')
plt.xlabel('Year')
plt.ylabel('Total Suicides')
```

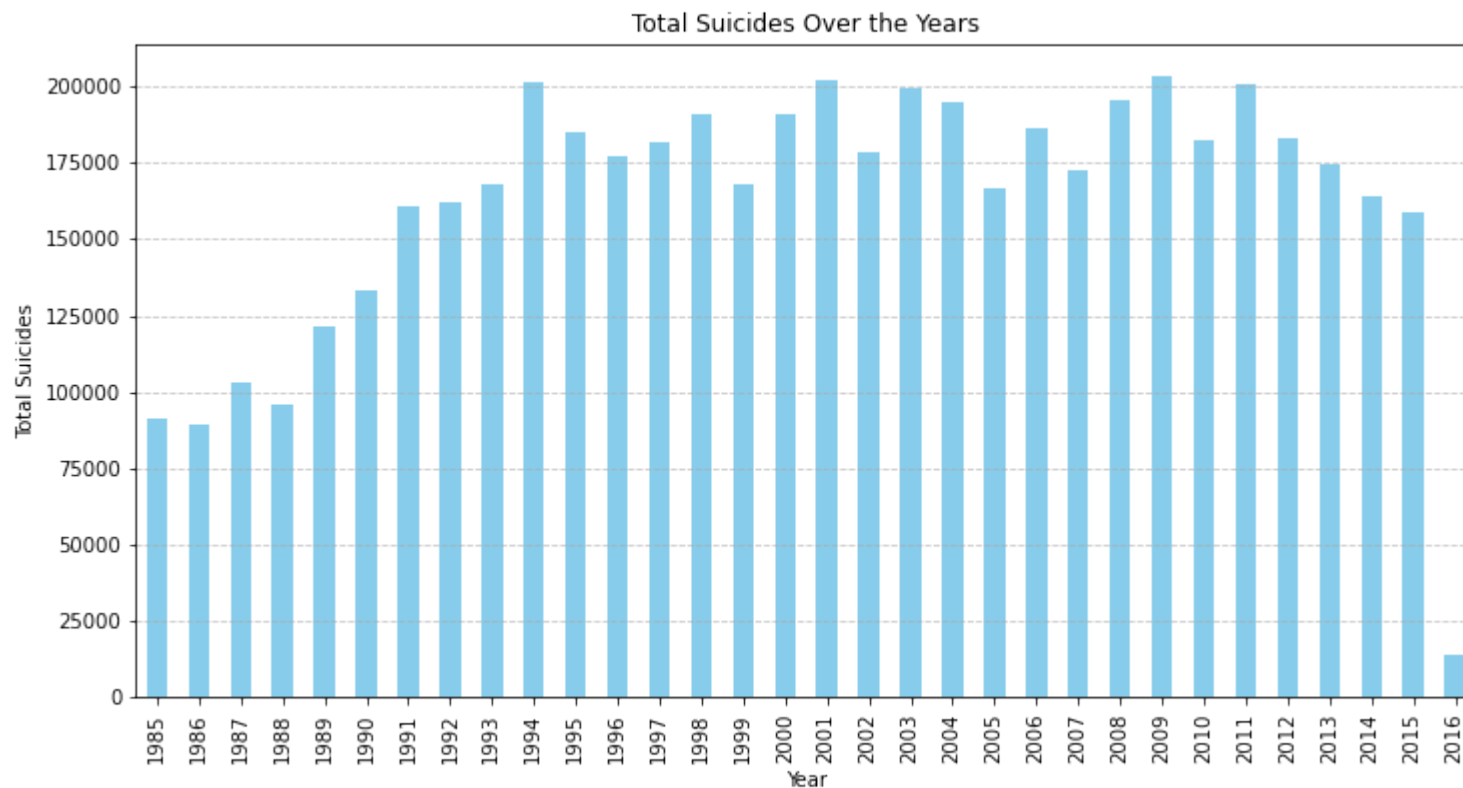


```
plt.grid(True)
plt.show()
```



```
In [54]: # Group the data by year and calculate the total suicides for each year
total_suicides_by_year = new_who.groupby('year')['suicides_no'].sum()

# Create a bar plot to visualize total suicides over the years
plt.figure(figsize=(12, 6))
total_suicides_by_year.plot(kind='bar', color='skyblue')
plt.title('Total Suicides Over the Years')
plt.xlabel('Year')
plt.ylabel('Total Suicides')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```




My Finding 2009 and 1994 are the years that went above 200,000 mark, while 2003 just fall on the mark. 1998, 200, 2004, 2008, 2003 fall in between 185,000 and 195,000. 2016 is the year with the lowest dead rate

Compare suicides by gender over years

```
In [55]: # Group the data by year and gender, and calculate the total suicides for each group
suicides_by_gender = new_who.groupby(['year', 'sex'])['suicides_no'].sum().reset_index()

# Create a line plot to compare suicides by gender over the years
plt.figure(figsize=(12, 6))
sns.lineplot(data=suicides_by_gender, x='year', y='suicides_no', hue='sex')
plt.title('Suicides by Gender Over the Years')
plt.xlabel('Year')
plt.ylabel('Total Suicides')
plt.grid(True)
plt.legend(title='Gender')
plt.show()
```

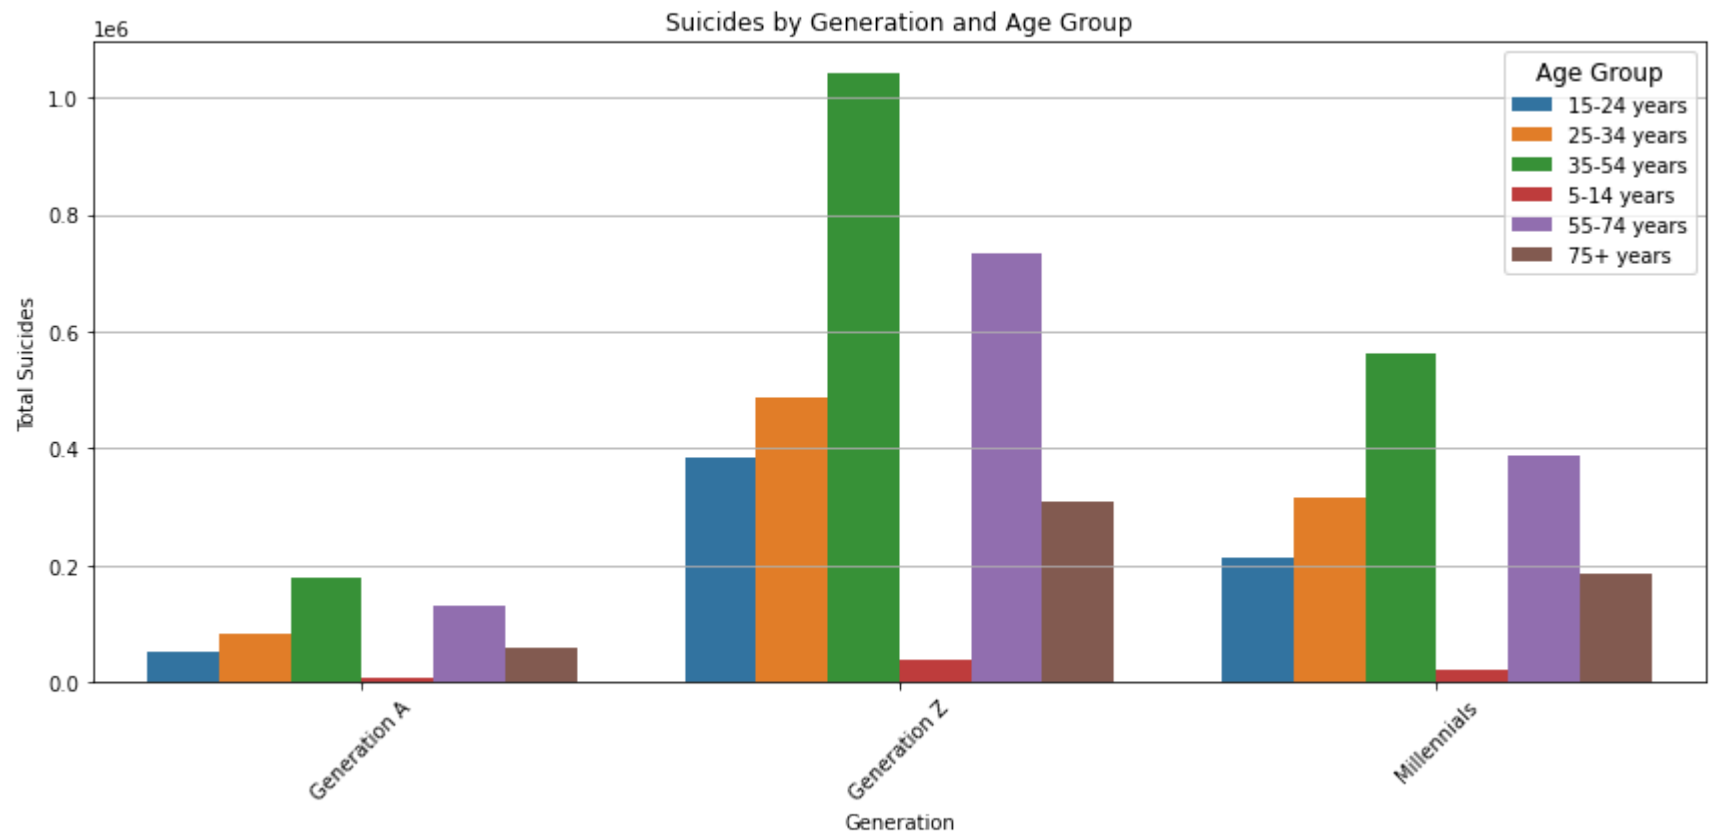


From the graph above it shows that male gender have a high rate in suicides and a sharp decline from 2015

Calculate and Visualise suicides on generation and on age group.

```
In [57]: # Group the data by generation and age group and calculate the total suicides for each group
suicides_by_generation_age = new_who.groupby(['generation', 'age'])['suicides_no'].sum().reset_index()

# Create a bar plot to visualize suicides by generation and age group
plt.figure(figsize=(12, 6))
sns.barplot(data=suicides_by_generation_age, x='generation', y='suicides_no', hue='age')
plt.title('Suicides by Generation and Age Group')
plt.xlabel('Generation')
plt.ylabel('Total Suicides')
plt.xticks(rotation=45)
plt.grid(axis='y')
plt.legend(title='Age Group', title_fontsize='12')
plt.tight_layout()
plt.show()
```



From all the generation A, Z and Millenials, we can see age 35-54,55-74 and 25-34 years, they are the genraation with the hight suicides rate.

## Section B

### Task Two, Car Sale on AA

This is to demostrate data acquiotion (Scraping, cleaning and exploration, However scraping could noot be done due to security for the website



```
In [65]: # importaing the data set  
cars=pd.read_csv('used_cars.csv')
```

```
In [66]: cars.head()
```

Out[66]:

	Sale title	Year	Body type	Number of seats	Number of reviews	Location	Distance	Mileage	Co2 emissions	Colour	Rating	price	Fuel type	Transmission	Engin siz
0	Toyota Yaris	2020	Hatchback	5	155	Steven Eagell Toyota Bedford	0	18,887	92 g/km	White	4.50	18890.00	Hybrid	Automatic	1.5
1	Toyota Yaris	2020	Hatchback	5	155	Steven Eagell Toyota Bedford	0	13,350	92 g/km	Red	4.50	19190.00	Hybrid	Automatic	1.5
2	Toyota Yaris	2020	Hatchback	5	155	Steven Eagell Toyota Bedford	0	13,552	92 g/km	White	4.50	19990.00	Hybrid	Automatic	1.5
3	Toyota C-HR	2020	Suv	5	1	Steven Eagell Toyota Bedford	0	55,220	86 g/km	Grey	2.50	20290.00	Hybrid	Automatic	1.8
4	Toyota Yaris	2021	Hatchback	5	155	Steven Eagell Toyota Bedford	0	14,156	92 g/km	Grey	4.50	20290.00	Hybrid	Automatic	1.5

## Checking the data types

In [67]: `cars.dtypes`

```
Out[67]: Sale title      object
         Year          int64
         Body type     object
         Number of seats int64
         Number of reviews int64
         Location      object
         Distance      int64
         Mileage        object
         Co2 emissions  object
         Colour         object
         Rating         float64
         price          float64
         Fuel type      object
         Transmission   object
         Engine size    object
         dtype: object
```

## Checking for NaN

```
In [68]: cars.isnull().sum()
```

```
Out[68]: Sale title      0
         Year          0
         Body type     0
         Number of seats 0
         Number of reviews 0
         Location      0
         Distance      0
         Mileage        0
         Co2 emissions  0
         Colour         0
         Rating         0
         price          0
         Fuel type      0
         Transmission   0
         Engine size    0
         dtype: int64
```

```
In [69]: # Read the dataset into a DataFrame (replace 'your_dataset.csv' with your file)
cars = pd.read_csv('used_cars.csv')

# Check for missing values in each column
missing_values = cars.isnull().sum()
```

```
# Check for unique values in each column
unique_values = cars.nunique()

# Check for data types of each column
data_types = cars.dtypes

# Combine the information into a summary DataFrame
summary_cars = pd.DataFrame({
    'Missing Values': missing_values,
    'Unique Values': unique_values,
    'Data Types': data_types
})

# Display columns with missing values
problematic_columns = summary_cars[summary_cars['Missing Values'] > 0]
print("Columns with missing values:")
print(problematic_columns)

# Display columns with too many unique values (potential categorical columns)
categorical_columns = summary_cars[summary_cars['Unique Values'] > len(cars) * 0.9]
print("\nPotential categorical columns:")
print(categorical_columns)

# Display columns with data type 'object' (usually text or categorical data)
text_columns = summary_cars[summary_cars['Data Types'] == 'object']
print("\nText (object) columns:")
print(text_columns)
```



Columns with missing values:

Empty DataFrame

Columns: [Missing Values, Unique Values, Data Types]

Index: []

Potential categorical columns:

Empty DataFrame

Columns: [Missing Values, Unique Values, Data Types]

Index: []

Text (object) columns:

	Missing Values	Unique Values	Data Types
Sale title	0	338	object
Body type	0	16	object
Location	0	30	object
Mileage	0	605	object
Co2 emissions	0	197	object
Colour	0	47	object
Fuel type	0	9	object
Transmission	0	5	object
Engine size	0	46	object

In [ ]: